

## Vaani Banking Voice Assistant

Round 2 Submission Document

Project Theme: AI Voice Assistant for Financial Operations

Repository: [Vaani Banking Voice Assistant : GitHub link](#)

Demo Video: [Demo Video link YouTube](#)

Submission Date: 23 Nov 2025

### Executive Summary

Vaani is an AI-powered voice banking assistant that enables customers to perform secure financial operations through natural conversational interactions. Built with a modern microservices architecture, Vaani supports multilingual operations (English & Hindi), implements comprehensive security guardrails, and provides seamless voice-first banking experiences compliant with RBI guidelines.

### Key Achievements:

- Full-stack voice-enabled banking platform with biometric authentication
- Multi-agent AI system with LangGraph orchestration
- Bilingual RAG system (English & Hindi) for loan and investment information
- Hello UPI payment system with RBI complianceComprehensive security guardrails (content moderation, PII detection, prompt injection protection)
- Production-ready architecture with observability and scalability

### 1. Technology Stack

#### 1.1 Frontend Layer

Framework & Build Tools:

- React 19 - Modern UI framework with hooks and context API

- Vite - Fast build tool and development server
- React Router v6 - Client-side routing
- Axios - HTTP client for API communication

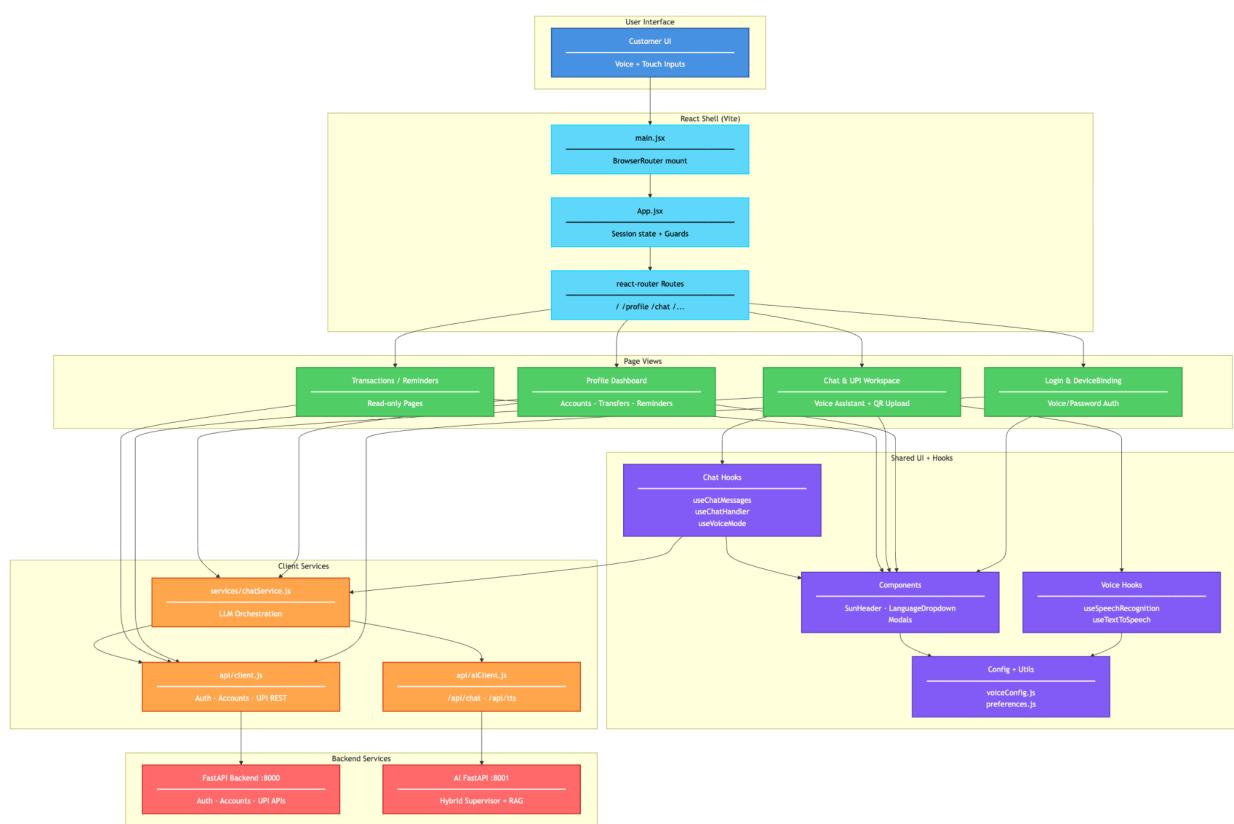
### Voice Technologies:

- Web Speech API - Browser-native Speech-to-Text (STT) and Text-to-Speech (TTS)
- Voice Recognition - Real-time voice input processing
- Voice Synthesis - Natural language audio output

### State Management:

- Context API - Global state management
- Local Storage - Session persistence and user preferences

### Frontend Architecture Diagram ([Link to source file](#))



## 1.2 Backend API Layer

### Framework & Runtime:

- FastAPI - High-performance async web framework
- Python 3.11+ - Modern Python with type hints

- Uvicorn - ASGI server for production deployment

## Database & ORM:

- SQLAlchemy 2.0 - Modern ORM with async support
  - SQLite (Development) - Lightweight database
  - PostgreSQL Ready - Production database configuration

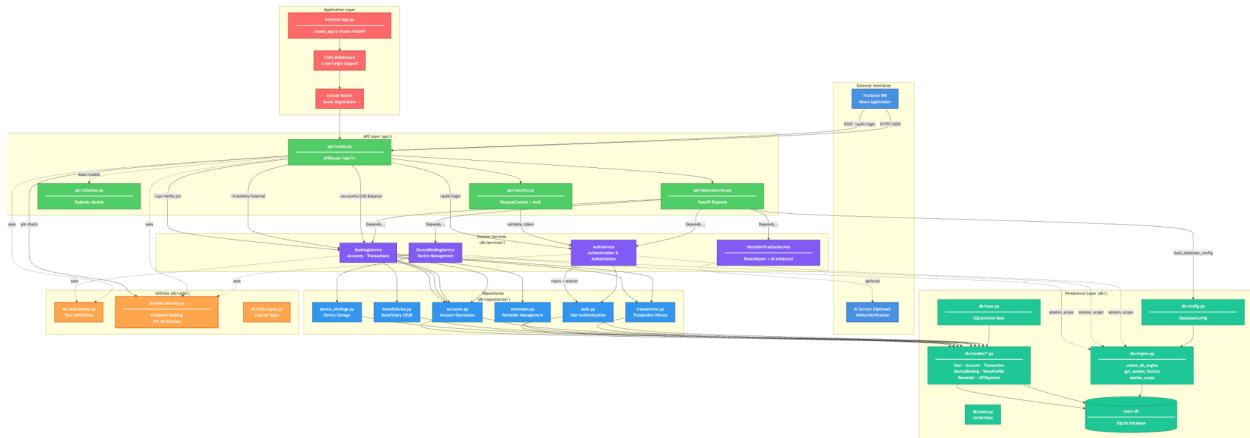
## Authentication & Security:

- python-jose - JWT token generation and validation (for production)
  - bcrypt - Password hashing (cost factor: 12)
  - Resemblyzer - Voice biometric verification
  - Pydantic v2 - Request/response validation

## API Features:

- RESTful API design
  - OpenAPI/Swagger documentation
  - CORS configuration
  - Request context tracking
  - Structured error handling

## Backend API Architecture ([link to source file](#))



### 1.3 AI Backend Layer

## AI Framework & Orchestration:

- LangGraph - Multi-agent orchestration and state management
  - LangChain - LLM integration and tool calling
  - LangSmith - Observability and tracing (optional)

## LLM Providers:

- Ollama (Primary) - Local LLM inference

- Qwen 2.5 7B - Comprehensive responses, multilingual support
- Llama 3.2 3B - Fast classification and voice mode
- OpenAI (Alternative) - Cloud-based LLM (configurable)

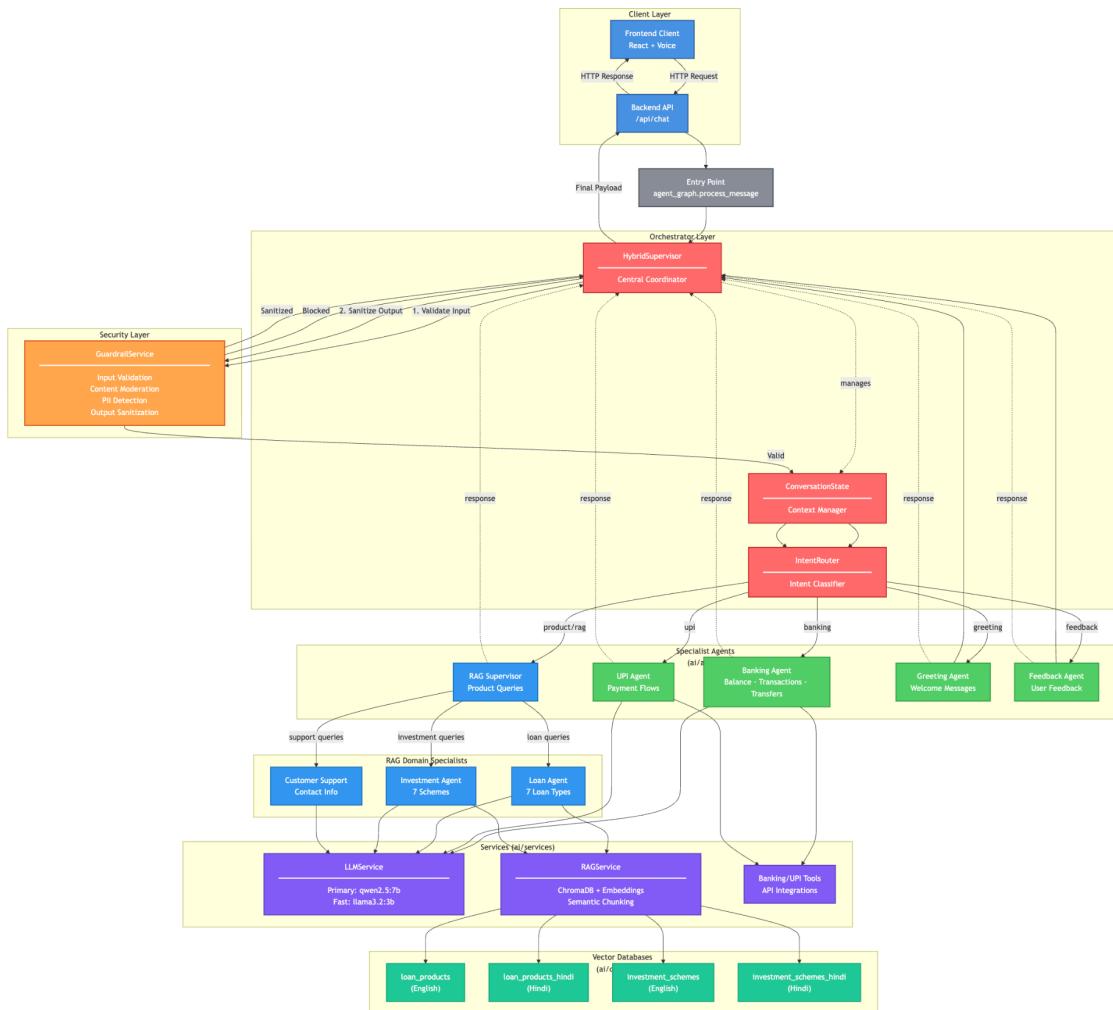
#### Vector Database & RAG:

- ChromaDB - Vector database for document embeddings
- sentence-transformers/all-MiniLM-L6-v2 - Multilingual embeddings (384 dimensions)
- Semantic Chunking - Intelligent document segmentation

#### AI Services:

- LLMService - Unified LLM provider abstraction
- RAGService - Retrieval-Augmented Generation with caching
- GuardrailService - Input/output safety checks
- Azure TTS (Optional) - Cloud-based text-to-speech

#### AI Architecture Diagram ([Link to source file](#))



## 1.4 Development & Deployment Tools

Package Management:

- uv - Fast Python package manager
- npm - Node.js package manager

Version Control:

- Git - Source code versioning
- GitHub - Repository hosting

Deployment (for production):

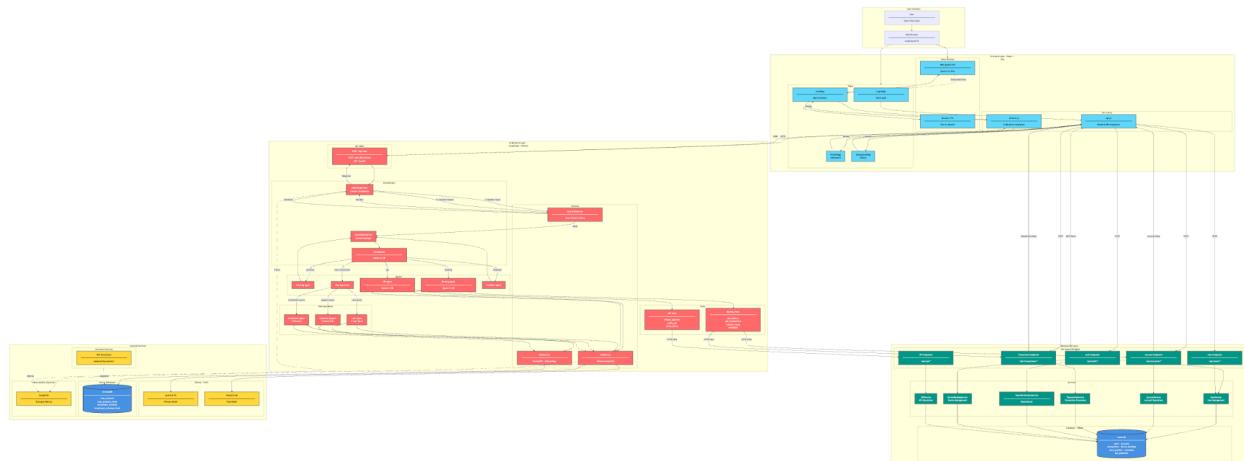
- Vercel - Frontend and backend deployment
- Docker (Ready) - Containerization support
- Kubernetes (Ready) - Orchestration support

Monitoring & Observability:

- structlog - Structured logging
- LangSmith - AI tracing and monitoring
- Python logging - Application logging

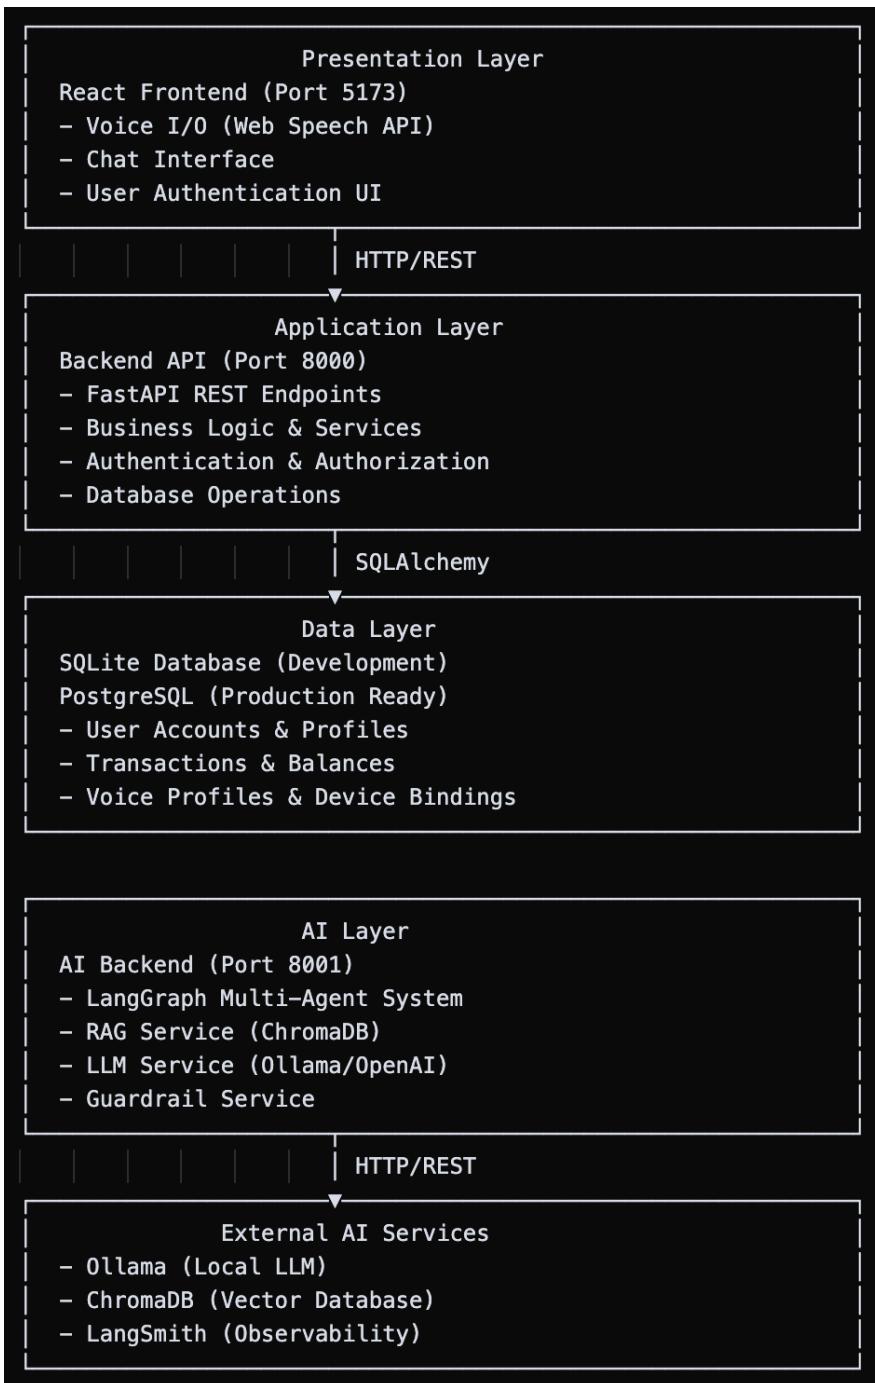
## 2. System Architecture

Technology Stack Overview ([Link to source file](#))



## 2.1 Three-Tier Microservices Architecture

Vaani follows a modern microservices architecture with clear separation of concerns:



## 2.2 Request Flow Architecture

### 2.2.1 Voice Login Flow

...

User speaks enrollment phrase



Frontend: Web Speech API transcribes



POST /api/auth/voice-login

{

```
username: "user123",
voice_sample: <audio_blob>,
device_info: {...}
```

}



Backend: Resemblyzer generates voice embedding



Backend: Compares with stored voice profile (cosine similarity)



Backend: Verifies device binding



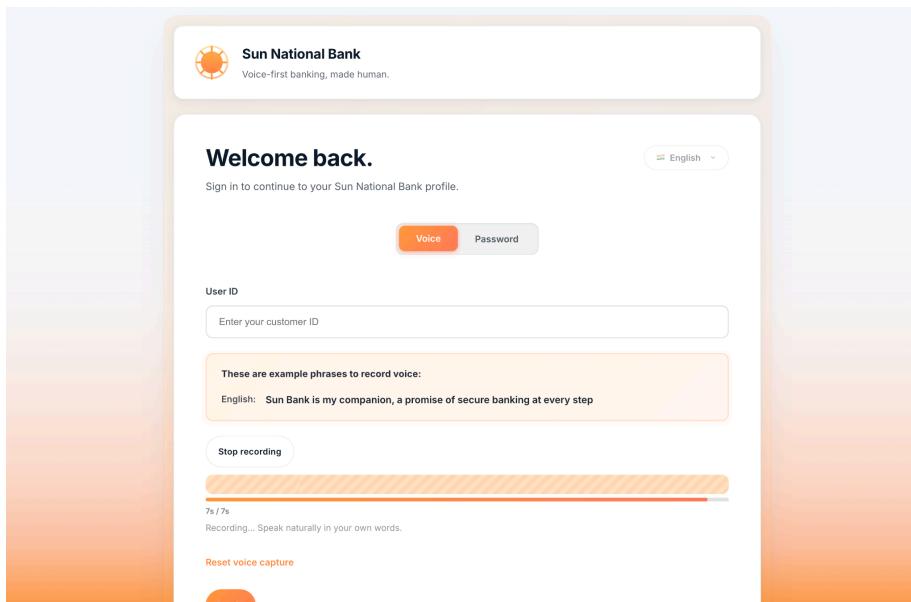
Backend: Returns JWT token (24h expiry)

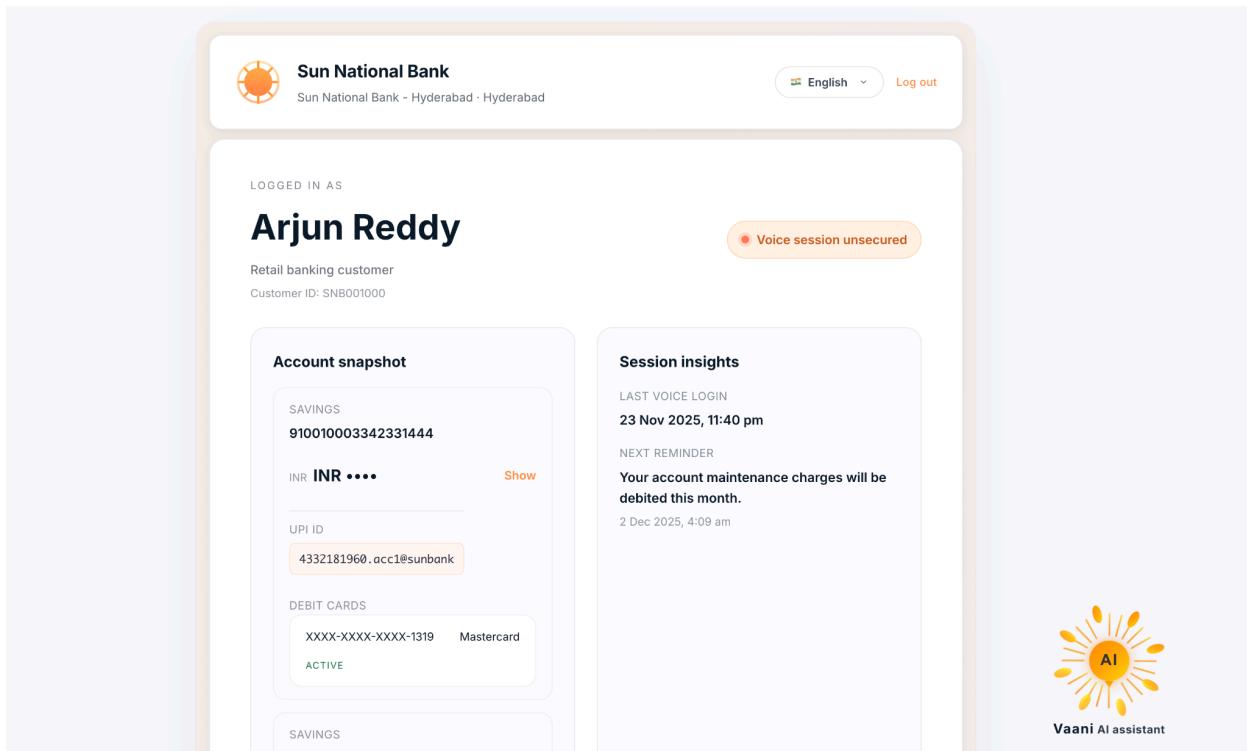


Frontend: Stores token, redirects to chat

...

### Voice Login Flow: Frontend login page screenshot





## 2.2.2 Banking Query Flow (Balance Check)

...

User: "What's my account balance?"



Frontend: aiClient.sendChatMessage()  
→ POST http://localhost:8001/api/chat



AI Backend: HybridSupervisor receives message



HybridSupervisor:  
1. Input Guardrails (content moderation, PII detection)  
2. IntentRouter classifies as "banking\_operation"  
3. Routes to Banking Agent



Banking Agent (Qwen 2.5 7B):

1. Decides to use get\_account\_balance tool
2. Tool calls Backend API:  
GET http://localhost:8000/api/accounts/{id}



Backend API:

1. Validates JWT token
2. Queries SQLite database

3. Returns: { balance: 50000.00, currency: "INR" }

↓

Banking Agent:

1. Formats response: "Your account balance is ₹50,000.00"

2. Output Guardrails (language consistency, safety check)

↓

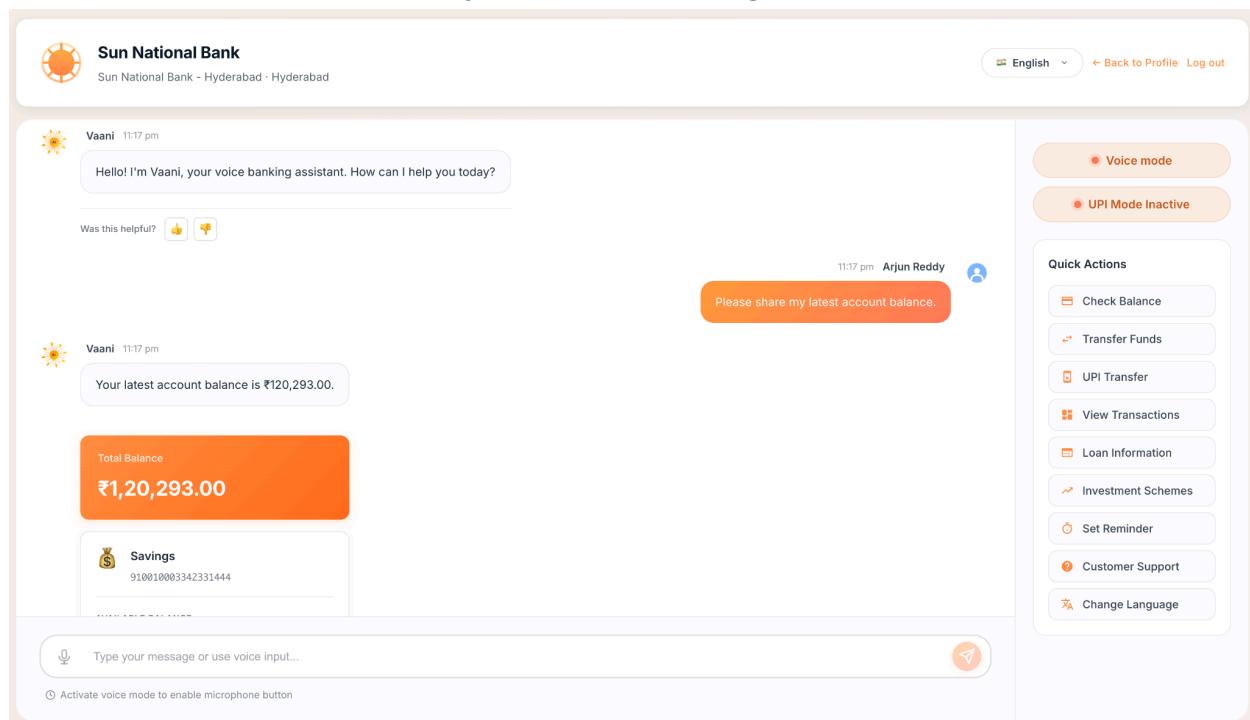
AI Backend returns response

↓

Frontend displays message + speaks via TTS

...

## Chat Interface with Balance Query : Frontend chat page screenshot



## 2.2.3 RAG Query Flow (Loan Information)

...

User: "Tell me about home loans" / "होम लोन के बारे में बताओ"



AI Backend: Intent classified as "loan\_inquiry"



RAG Supervisor: Routes to Loan Agent



Loan Agent:

1. Detects language (English or Hindi)
2. Queries appropriate vector DB:
  - loan\_products (English)
  - loan\_products\_hindi (Hindi)
3. Retrieves relevant PDF chunks (semantic search)
4. Sends context + query to LLM (Qwen 2.5 7B)



LLM generates structured response with loan card



Frontend displays loan information card

...

## Loan Information Card Display: Frontend chat page showing loan card

The screenshot shows a mobile application interface for Sun National Bank. At the top, there's a header with the bank's logo and name, language selection (Hindi), and a button to log out. Below the header, a message from 'Vaani' at 11:19 pm is shown: "होम लोन के लिए हमारे पास 6 विकल्प उपलब्ध हैं। कृपया अपनी आवश्यकता के अनुसार चुनें।". A reply from 'Arjun Reddy' follows: "home loan के बारे में बताएं". To the right, there's a sidebar titled "त्वरित क्रियाएँ" (Quick Actions) with options like "खाता बैलेस देखें", "UPI से पैसा ट्रांसफर करें", etc. The main area displays a "उपलब्ध ऋण प्रकार" (Available Loan Types) card with four items: "होम पर्फेज लोन" (Home Perfex Loan), "होम कंस्ट्रक्शन लोन" (Home Construction Loan), "प्लॉट कंस्ट्रक्शन लोन" (Plot Construction Loan), and "होम एक्स्टेंशन लोन" (Home Extension Loan). At the bottom, there are buttons for "अपना जादेश राइप करें या आदाज का उपयोग करें..." and "गाइडलाइन बदल सकते हैं बदलें".

## 2.2.4 UPI Payment Flow

...

User: "Hello Vaani, send ₹500 to John via UPI"



AI Backend: Intent classified as "upi\_payment"



UPI Agent (Multi-step flow):

Step 1: Confirm recipient and amount

→ "Confirming: Send ₹500 to John?"

User: "Yes"

Step 2: Request UPI PIN

→ "Please enter your UPI PIN"

→ structured\_data: { action: "collect\_upi\_pin" }

Frontend: Shows PIN input modal

User enters PIN

Step 3: Verify PIN

→ Tool calls: POST /api/upi/verify-pin

Step 4: Process payment

→ Tool calls: POST /api/upi/send-payment

Backend:

1. Verifies PIN
2. Creates debit/credit transactions
3. Updates balances
4. Returns confirmation

Step 5: Confirmation

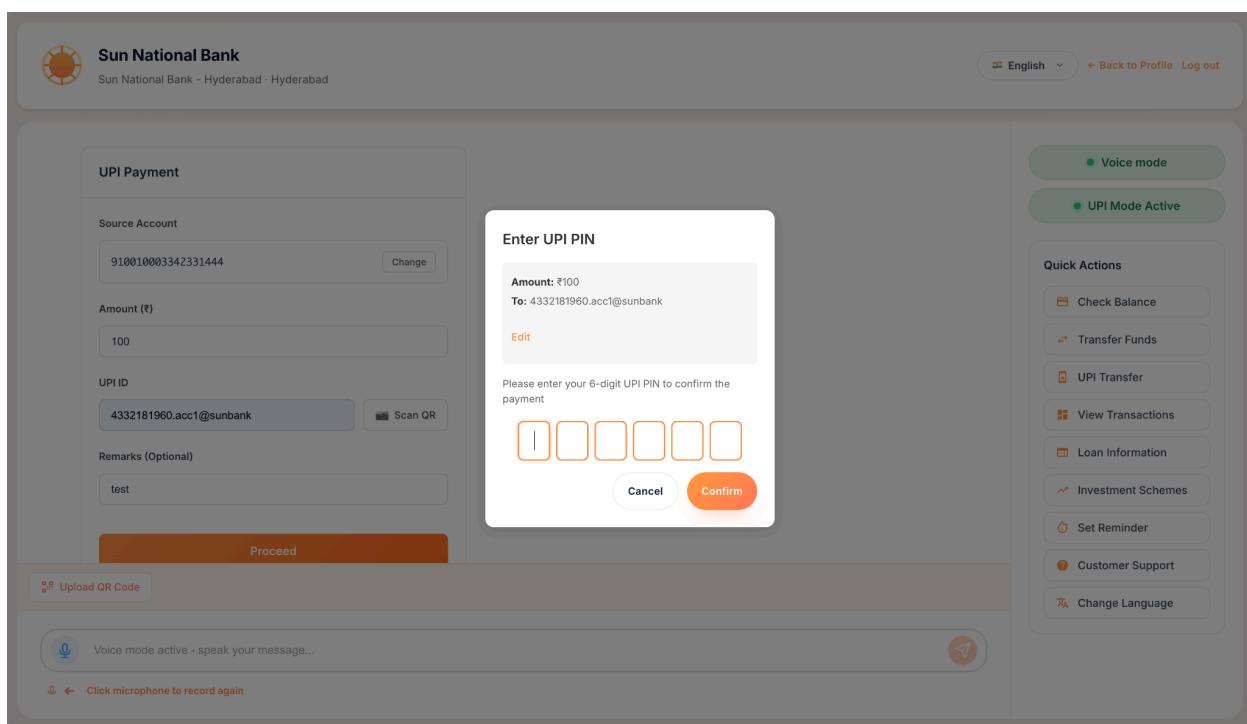
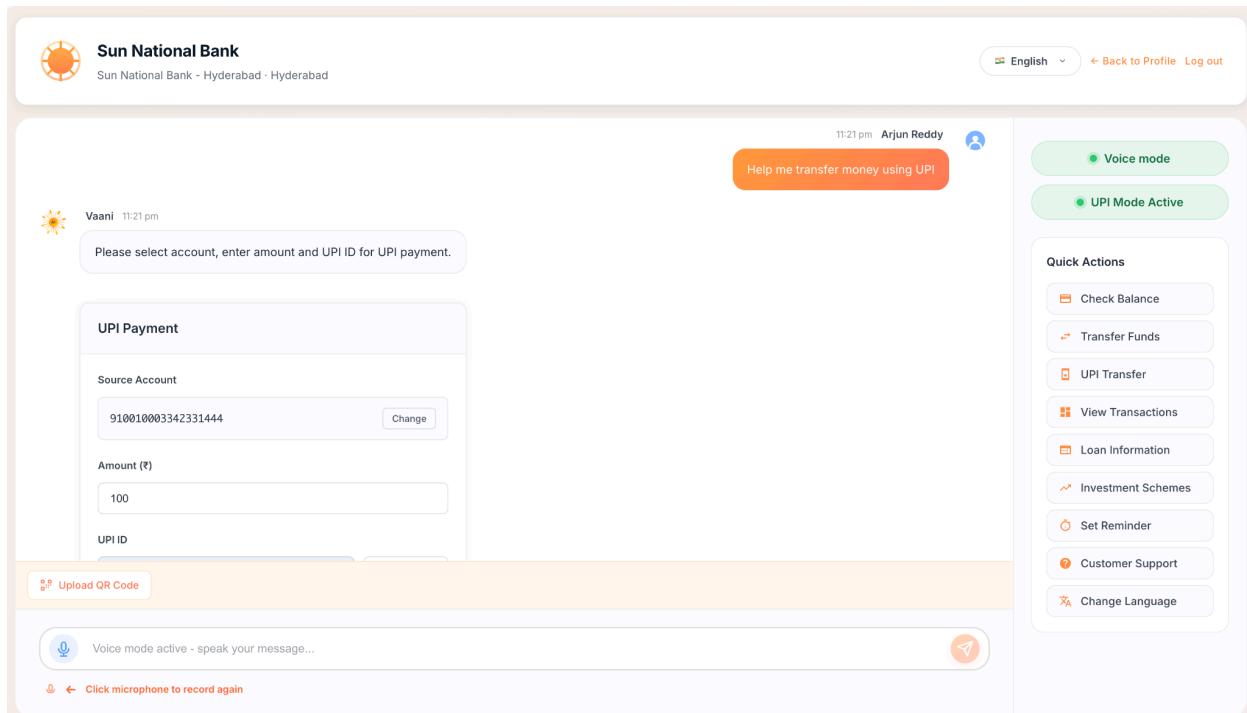
→ "Payment successful! ₹500 sent to John."

→ structured\_data: { payment\_id: "TXN123", status: "success" }

...

...

## UPI Payment Flow : Frontend showing UPI payment modal, PIN entry, success message



## Other banking operations:

### Set reminder

The screenshot shows the Sun National Bank mobile application interface. At the top, there is a header with the bank's logo, name, language selection (English), and navigation links (Back to Profile, Log out). The main content area has a title "Set a reminder" and a sub-instruction: "You can create payment reminders here. Fill in the form below with date, time, account, and message. Reminders can be sent via voice, SMS, or push notifications." Below this, there is a form for setting a reminder. The fields include:

- Reminder type:** Bill payment
- Message:** e.g. Pay electricity bill
- Remind me on:** 24/11/2025, 06:11 PM
- Linked account (optional):** (empty)

At the bottom of the form is a text input field with placeholder "Type your message or use voice input..." and a send button.

On the right side of the screen, there is a sidebar titled "Quick Actions" containing the following options:

- Check Balance
- Transfer Funds
- UPI Transfer
- View Transactions
- Loan Information
- Investment Schemes
- Set Reminder
- Customer Support
- Change Language

There are also two status indicators at the top right: "Voice session unsecured" and "UPI Mode Inactive".

### Get Investment details

The screenshot shows the Sun National Bank mobile application interface. At the top, there is a header with the bank's logo, name, language selection (English), and navigation links (Back to Profile, Log out). The main content area has a title "Available Investment Schemes" and a sub-instruction: "Here are the available investment schemes. Click or speak any scheme for detailed information:". Below this, there is a list of investment schemes:

Scheme	Description	Action
PPF	Long-term tax-saving scheme	→
NPS	Market-linked retirement scheme	→
Sukanya Samridhi Yojana	Girl child savings scheme	→
ELSS	Tax-saving mutual funds	→

At the bottom of the screen is a text input field with placeholder "Type your message or use voice input..." and a send button.

On the right side of the screen, there is a sidebar titled "Quick Actions" containing the following options:

- Check Balance
- Transfer Funds
- UPI Transfer
- View Transactions
- Loan Information
- Investment Schemes
- Set Reminder
- Customer Support
- Change Language

There are also two status indicators at the top right: "Voice session unsecured" and "UPI Mode Inactive".

## Get Account Statement:

Sun National Bank  
Sun National Bank - Hyderabad · Hyderabad

English Back to Profile Log out

Vaani 11:51 pm Arjun Reddy

Show me seven days account statement

Voice session unsecured  
UPI Mode Inactive

**Download Statement**  
Select an account and period to download your statement.

Select account: savings • 910010003342331444

Select period: Last 7 days (highlighted), Last 30 days, Last 3 months, Last 6 months, Last 12 months, Custom range

Download Statement

Was this helpful?

Type your message or use voice input...

Activate voice mode to enable microphone button

Quick Actions

- Check Balance
- Transfer Funds
- UPI Transfer
- View Transactions
- Loan Information
- Investment Schemes
- Set Reminder
- Customer Support
- Change Language

localhost:5173/chat

snb\_statement\_31444\_20251116\_to\_20251123

snb_statement_31444_20251116_to_20251123							
Sun National Bank	Account Statement						
Account Holder	Arjun Reddy						
Account Number	910010003342331444						
Account Type	SAVINGS						
Statement Period	2025-11-16 to 2025-11-23						
Opening Balance (INR)	132556.00						
Closing Balance (INR)	120293.00						
Generated On	–						
Transaction Date	Value Date	Description	Reference No.	Debit (INR)	Credit (INR)	Balance (INR)	Status
17 Nov 2025	17 Nov 2025	Movie ticket booking	66f9474f-ccf	12263.00		120293.00	settled

Sun National Bank  
Sun National Bank - Hyderabad · Hyderabad

Download Statement  
Select an account and period to download your statement.

Select account: savings • 910010003342331444

Select period: Last 7 days (highlighted), Last 30 days, Last 3 months, Last 6 months, Last 12 months, Custom range

Statement prepared. Download started.

Download Statement

Was this helpful?

Vaani 11:52 pm

Account statement downloaded successfully. Account: 31444 (2025-11-16 to 2025-11-23)

Type your message or use voice input...

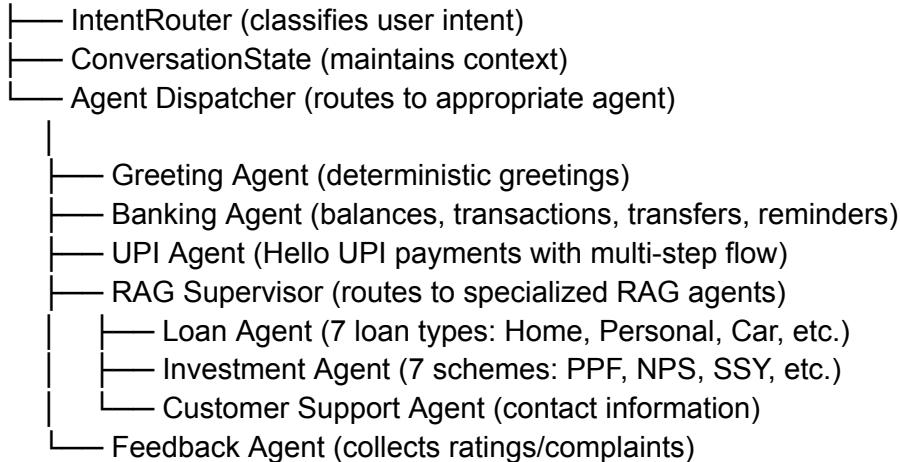
Activate voice mode to enable microphone button

## 2.3 Multi-Agent Architecture

Vaani uses a Hybrid Supervisor Pattern with specialized agents:

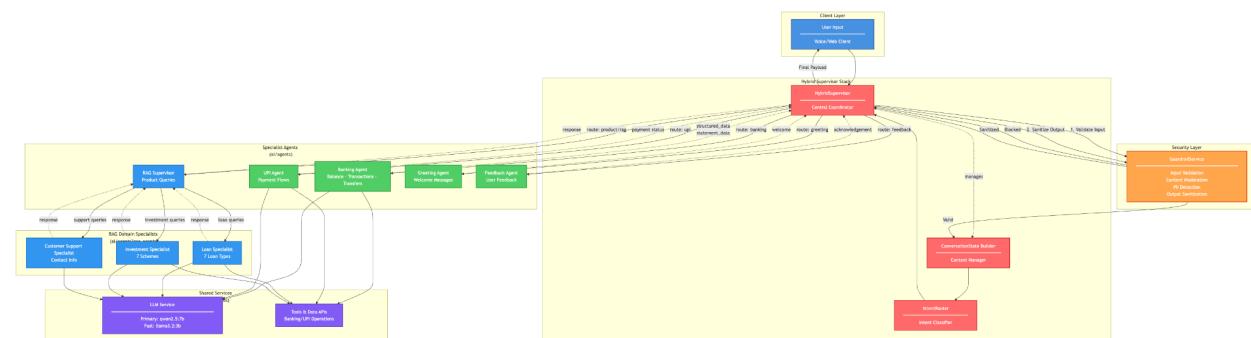
...

### HybridSupervisor



...

### Agent Architecture Diagram: Agent hierarchy, routing logic, tool calling ([Link to source file](#))



## 3. Data Model & Storage

### 3.1 Database Schema

Primary Database: SQLite (Development) / PostgreSQL (Production Ready)

#### Core Tables

### Users Table

```
```sql
CREATE TABLE users (
    id VARCHAR PRIMARY KEY,
    username VARCHAR UNIQUE NOT NULL,
    password_hash VARCHAR NOT NULL,
    first_name VARCHAR,
    last_name VARCHAR,
    email VARCHAR,
    phone_number VARCHAR,
    upi_id VARCHAR UNIQUE,
    created_at TIMESTAMP,
    updated_at TIMESTAMP
);
````
```

### Accounts Table

```
```sql
CREATE TABLE accounts (
    id VARCHAR PRIMARY KEY,
    user_id VARCHAR REFERENCES users(id),
    account_number VARCHAR UNIQUE NOT NULL,
    account_type VARCHAR, -- SAVINGS, CHECKING, etc.
    balance DECIMAL(15, 2) DEFAULT 0.00,
    currency VARCHAR DEFAULT 'INR',
    created_at TIMESTAMP,
    updated_at TIMESTAMP
);
````
```

### Transactions Table

```
```sql
CREATE TABLE transactions (
    id VARCHAR PRIMARY KEY,
    account_id VARCHAR REFERENCES accounts(id),
    transaction_type VARCHAR, -- DEBIT, CREDIT
    amount DECIMAL(15, 2) NOT NULL,
    channel VARCHAR, -- UPI, TRANSFER, ATM, etc.
    reference_id VARCHAR,
    description TEXT,
    created_at TIMESTAMP
);
````
```

### Voice Profiles Table

```
```sql
CREATE TABLE voice_profiles (
    id VARCHAR PRIMARY KEY,
    user_id VARCHAR REFERENCES users(id),
    voice_embedding BLOB, -- Resemblyzer embedding (256 floats)
    enrollment_phrase TEXT,
    created_at TIMESTAMP,
    updated_at TIMESTAMP
);
````
```

### Device Bindings Table

```
```sql
CREATE TABLE device_bindings (
    id VARCHAR PRIMARY KEY,
    user_id VARCHAR REFERENCES users(id),
    device_identifier VARCHAR NOT NULL,
    device_fingerprint VARCHAR,
    device_label VARCHAR,
    platform VARCHAR,
    is_trusted BOOLEAN DEFAULT FALSE,
    created_at TIMESTAMP,
    last_used_at TIMESTAMP
);
````
```

### Reminders Table

```
```sql
CREATE TABLE reminders (
    id VARCHAR PRIMARY KEY,
    user_id VARCHAR REFERENCES users(id),
    title VARCHAR NOT NULL,
    description TEXT,
    reminder_type VARCHAR, -- PAYMENT, BILL, etc.
    due_date DATE,
    status VARCHAR, -- PENDING, COMPLETED, CANCELLED
    created_at TIMESTAMP
);
````
```

### UPI Payments Table

```
```sql
CREATE TABLE upi_payments (
```

```
    id VARCHAR PRIMARY KEY,  
    user_id VARCHAR REFERENCES users(id),  
    source_account_id VARCHAR REFERENCES accounts(id),  
    destination_account_id VARCHAR REFERENCES accounts(id),  
    amount DECIMAL(15, 2) NOT NULL,  
    upi_reference_id VARCHAR UNIQUE,  
    status VARCHAR, -- PENDING, SUCCESS, FAILED  
    created_at TIMESTAMP  
);  
...
```

### 3.2 Vector Database (RAG System)

Vector Database: ChromaDB

Collections

1. Loan Products (English)

- Collection: `loan\_products`
- Documents: 7 PDF files (Home, Personal, Car, Education, Business, Gold, Agriculture loans)
- Embeddings: `sentence-transformers/all-MiniLM-L6-v2` (384 dimensions)
- Chunks: ~350 semantic chunks with metadata

Note: Similarly for other documents on investments and loans

### 3.3 Data Storage Strategy

**Structured Data:**

- SQLite (Development) - File-based, zero configuration
- PostgreSQL (Production) - Managed service (AWS RDS, Supabase)

**Vector Data:**

- ChromaDB - Local persistence with `persist\_directory`
- Future: Managed ChromaDB or Pinecone/Weaviate for production

**Voice Data:**

- Embeddings Only - Stored as BLOB (256 floats per embedding)
- Raw Audio - Not stored (privacy-first approach)
- Future: Encrypted storage for production

**File Storage:**

- PDF Documents - Local filesystem (`backend/documents/`)
- Future: S3/blob storage for production

## **4. AI / ML / Automation Components**

### **4.1 Multi-Agent System (LangGraph)**

Vaani implements a sophisticated multi-agent system using LangGraph for orchestration:

#### **4.1.1 Intent Classification**

Model: Llama 3.2 3B (Fast Model)

Purpose: Quick intent classification for routing

Accuracy: ~90%+ intent detection

#### **4.1.2 Banking Agent**

Model: Qwen 2.5 7B

#### **4.1.3 UPI Agent**

Model: Qwen 2.5 7B

### **4.1.4 RAG System (Retrieval-Augmented Generation)**

Architecture:

- Supervisor Pattern: RAG Supervisor routes to specialized agents
- Loan Agent: Handles 7 loan types with bilingual support
- Investment Agent: Handles 7 investment schemes with bilingual support
- Customer Support Agent: Provides contact information

RAG Process:

1. Query Processing: User query in English or Hindi
2. Language Detection: Automatic language identification
3. Vector Search: Semantic similarity search in appropriate database
4. Context Retrieval: Top-k relevant chunks (k=3-5)
5. LLM Generation: Qwen 2.5 7B generates response with context
6. Structured Output: Returns loan/investment card with structured data

Performance:

- Retrieval Latency: 50-200ms
- LLM Generation: 500-1000ms
- Total RAG Query: < 3 seconds
- Cache Hit Rate: ~40% (120s TTL)

## **4.2 Voice Biometrics**

Technology: Resemblyzer

Purpose: Voice-based authentication

Process:

1. Enrollment: User speaks enrollment phrase, embedding generated (256 floats)
2. Storage: Embedding stored in database (not raw audio)
3. Verification: New voice sample compared using cosine similarity
4. Threshold: 0.6 similarity score for authentication

Security Features:

- Device binding required for voice login
- Single trusted device at a time
- AI-enhanced verification (optional augmentation)

## **4.3 Guardrails & Safety**

Implementation: Open-source only, no external API dependencies

### **4.3.1 Input Guardrails**

Content Moderation:

- Toxic content detection (English & Hindi keywords)
- Profanity, threats, hate speech, harassment detection

PII Detection:

- Aadhaar numbers (12 digits)
- PAN cards (5 letters + 4 digits + 1 letter)
- Account numbers (9-18 digits)
- CVV/PIN (3-6 digits with keywords)
- Phone numbers, card numbers, IFSC codes

Prompt Injection Protection:

- Detects attempts to override system instructions
- Blocks commands like "ignore previous instructions"
- Supports English and Hindi injection patterns

Rate Limiting:

- Per-user limits: 30 requests/minute, 500 requests/hour
- In-memory tracking (Redis-ready for distributed)

### **4.3.2 Output Guardrails**

Language Consistency:

- Verifies response matches requested language
- Detects language mixing (max 30% mixing allowed)
- Ensures Hindi responses use Devanagari script

Response Safety:

- Checks AI output for toxic content
- Prevents PII leakage in responses
- Validates structured data format

#### **4.4 Multilingual Support**

Languages Supported:

- English (en-IN) - Full support
- Hindi (hi-IN) - Full support with Devanagari script

Implementation:

- Bilingual RAG Databases: Separate vector databases for English and Hindi
- Language Detection: Automatic detection from user query
- LLM Multilingual: Qwen 2.5 7B supports 100+ languages
- Embeddings: Multilingual model (all-MiniLM-L6-v2) works for both languages

Hindi Features:

- Native Hindi PDF documents with proper fonts
- Hindi loan and investment information
- Hindi conversational responses
- Cultural appropriateness (female gender for Vaani)

#### **4.5 Observability & Monitoring**

LangSmith Integration:

- Tracing: All LLM interactions traced
- Metrics: Token usage, latency, error rates
- Visualization: Trace waterfall, conversation flows
- Debugging: Agent decisions, tool calls, errors

Structured Logging:

- Backend: Python logging with request/response tracking
- AI Backend: structlog for structured events
- Events: LLM calls, tool executions, agent decisions, errors

Performance Metrics:

- Response time tracking

- Token usage monitoring
- Cache hit rates
- Error rate tracking

## 5. Security & Compliance

### 5.1 Authentication & Authorization

#### 5.1.1 Multi-Factor Authentication

Methods Supported:

1. Username/Password - Traditional authentication
2. Voice Biometric - Resemblyzer-based voice verification
3. Device Binding - Device fingerprint verification
4. OTP Validation - One-time password (future enhancement)

Voice Authentication Flow:

- User speaks enrollment phrase
- Voice embedding generated (256-dimensional vector)
- Embedding stored securely (not raw audio)
- Future logins: voice sample compared using cosine similarity
- Threshold: 0.6 similarity score
- Device binding required for security

#### 5.1.2 JWT Token Security

Implementation:

- Algorithm: HS256
- Expiry: 24 hours (configurable)
- Storage: Local storage (future: HTTP-only cookies)
- Validation: All protected routes validate JWT

Token Structure:

```
```json
{
  "user_id": "uuid",
  "username": "user123",
  "exp": 1234567890,
  "iat": 1234567890
}
```

```

## 5.2 Data Security

### 5.2.1 Password Security

- Hashing: bcrypt with cost factor 12
- Storage: Hashed passwords only (never plaintext)
- Validation: Secure comparison using bcrypt

### 5.2.2 Voice Data Security

- Storage: Embeddings only (256 floats), not raw audio
- Privacy: Raw voice samples never stored
- Encryption: Future: encrypted embeddings at rest

### 5.2.3 Database Security

- Development: SQLite with file permissions
- Production: PostgreSQL with encryption at rest
- Connection: TLS for database connections (production)

### 5.2.4 API Security

- CORS: Configured origins (development: permissive, production: strict)
- Input Validation: Pydantic schemas for all requests
- Rate Limiting: Guardrail-based rate limiting (30/min, 500/hour)
- HTTPS: TLS for all API calls (production)

## 5.3 RBI Compliance (UPI Payments)

### 5.3.1 Hello UPI Compliance

#### PIN Security:

- PIN must be entered manually (keyboard/screen)
- PIN never spoken aloud
- PIN masked during entry
- PIN cleared immediately after use
- PIN not stored anywhere

#### User Consent:

- Explicit consent before first UPI payment
- Terms and conditions displayed
- Consent recorded and stored
- Can be revoked

Transaction Verification:

- Amount, recipient, and account displayed before PIN
- User confirms before PIN entry
- Can cancel at any time
- Confirmation after success

Audit Trail:

- Every transaction has unique reference ID (UPI-YYYYMMDD-HHMMSS)
- UPI channel marked in transactions
- Complete transaction log
- Timestamps in IST

## 5.4 Guardrails & Content Safety

Comprehensive Guardrails:

- Input Guardrails: Content moderation, PII detection, prompt injection protection, rate limiting
- Output Guardrails: Language consistency, response safety, PII leakage prevention
- Open-Source Only: No external API dependencies
- Bilingual Support: English and Hindi guardrails

## 5.5 Privacy & Data Protection

Privacy-First Design:

- Voice samples not stored (embeddings only)
- PII detection prevents accidental sharing
- No tracking or analytics without consent
- User data encrypted at rest (production)

Data Minimization:

- Only necessary data collected
- Voice embeddings minimal (256 floats)
- Transaction data minimal (amount, parties, timestamp)

---

## 6. Scalability & Performance

### 6.1 Architecture Scalability

#### 6.1.1 Horizontal Scaling

Frontend:

- Static Files: Infinitely scalable via CDN
- No State: Can serve unlimited concurrent users

- Deployment: Vercel, Netlify, S3 + CloudFront

Backend API:

- Stateless Design: JWT tokens enable horizontal scaling
- Load Balancer: Multiple instances behind load balancer
- Database Pooling: Connection pooling for database access
- Deployment: Docker containers, Kubernetes orchestration

AI Backend:

- Stateless Agents: Agent execution is stateless
- Multiple Instances: Can scale with multiple instances
- Shared ChromaDB: Managed service for vector database
- LLM Scaling: GPU instances for Ollama or cloud LLM APIs

### 6.1.2 Caching Strategy

Current Implementation:

- RAG Context Cache: 120-second TTL for retrieved context
- In-Memory Caching: Per-instance caching
- Cache Hit Rate: ~40% for frequent queries

Future Enhancements:

- Redis: Distributed caching for multiple instances
- Query Caching: Cache frequent queries (balance checks)
- Session State Caching: User session state caching
- Vector Search Caching: Cache vector search results

## 6.2 Performance Metrics

### 6.2.1 Response Times

Target Performance:

- Voice Mode (Fast Model): < 1 second
- Standard Chat: < 2 seconds
- RAG Queries: < 3 seconds
- Complex Banking Ops: < 5 seconds

Current Performance (Apple M4 Max, 128GB RAM):

- Intent Classification (Llama 3.2 3B): 100-200ms
- Balance Query (Qwen 2.5 7B): 500-800ms
- General FAQ (Qwen 2.5 7B): 1-2s
- RAG Retrieval: 50-200ms
- LLM Generation: 500-1000ms
- UPI Payment: 1-1.5s

## 6.2.2 Optimization Strategies

Model Selection:

- Fast Model (Llama 3.2 3B): Used for voice mode and intent classification
- Comprehensive Model (Qwen 2.5 7B): Used for complex queries and RAG

RAG Optimization:

- Semantic Chunking: Intelligent document segmentation
- Metadata Filtering: Language and document type filtering
- Context Caching: 120-second TTL reduces redundant retrievals

Future Optimizations:

- Streaming Responses: First token in 200-400ms
- Parallel Tool Calls: Execute multiple tools concurrently
- Embedding Caching: Cache document embeddings
- Warm Start: Pre-load models for faster first request

## 6.3 Error Handling & Resilience

### 6.3.1 Retry Strategy

Implementation:

- Automatic retry (3 attempts)
- Exponential backoff (1s, 2s, 4s)
- Graceful degradation on failure

### 6.3.2 Circuit Breaker Pattern

Future Implementation:

- Monitor Backend API health
- Open circuit after 5 failures in 1 minute
- Use cached data or fallback
- Half-open after 30 seconds for testing

### 6.3.3 Graceful Degradation

Fallback Chain:

1. Ollama unavailable → Fallback to OpenAI API (if configured)
2. OpenAI unavailable → Use rule-based responses
3. No LLM available → Return helpful error message

## **6.4 Database Performance**

Current Setup:

- SQLite: File-based, suitable for development
- Connection Pooling: SQLAlchemy connection pooling
- Query Optimization: Indexed queries, efficient joins

Production Ready:

- PostgreSQL: Managed service (AWS RDS, Supabase)
- Read Replicas: For read-heavy workloads
- Connection Pooling: PgBouncer or SQLAlchemy pooling
- Query Optimization: Indexes, query analysis, optimization

## **7. Innovation Highlights**

### **7.1 Novel Features**

#### **7.1.1 Hybrid Supervisor Pattern**

- Innovation: Multi-level agent orchestration with supervisor routing
- Benefit: Specialized agents for different domains, better accuracy
- Differentiator: Not just a single LLM, but a coordinated multi-agent system

#### **7.1.2 Bilingual RAG System**

- Innovation: Separate vector databases for English and Hindi with same embedding model
- Benefit: Native language support without performance penalty
- Differentiator: True multilingual support, not just translation

#### **7.1.3 Voice-First Banking**

- Innovation: Complete voice-enabled banking operations
- Benefit: Hands-free banking, accessibility for all users
- Differentiator: Not just voice commands, but natural conversation

#### **7.1.4 Comprehensive Guardrails**

- Innovation: Open-source guardrails with no external API dependencies
- Benefit: Privacy-preserving, cost-effective, customizable
- Differentiator: Production-ready safety without vendor lock-in

#### **7.1.5 Hello UPI Implementation**

- Innovation: Voice-assisted UPI payments with RBI compliance
- Benefit: Secure, compliant, user-friendly payment experience
- Differentiator: Full UPI flow with voice, not just balance checks

## 7.2 Technical Differentiators

1. Multi-Agent Architecture: Not a single LLM, but coordinated agents
  2. Bilingual Native Support: True Hindi support, not translation
  3. Production-Ready Guardrails: Open-source, no vendor lock-in
  4. Voice Biometrics: Secure voice authentication
  5. RBI-Compliant UPI: Full payment flow with compliance
  6. Observability: LangSmith integration for debugging
  7. Scalable Architecture: Microservices, stateless design, horizontal scaling
- 
8. Business Impact & Outcomes

### 8.1 Core Banking Operations Coverage

Fully Implemented:

- Account balance checking
- Transaction history viewing
- Fund transfers between accounts
- Payment reminders (create, list, update)
- Loan information (7 types)
- Investment scheme information (7 schemes)
- UPI payments (Hello UPI)
- Customer support queries

### 8.2 Accessibility & Inclusion

Voice-First Design:

- Hands-free banking for all users
- Natural language understanding
- No complex menu navigation
- Multilingual support (English & Hindi)

User Benefits:

- Elderly users: Simple voice commands
- Visually impaired: Voice-only interface
- Non-tech-savvy users: Natural conversation
- Regional language users: Hindi support

### 8.3 Business Metrics & ROI

#### **Expected Outcomes:**

- Reduced Call Center Load: Voice assistant handles common queries
- 24/7 Availability: Always-on banking assistant
- Faster Transactions: Voice commands faster than app navigation
- Higher Engagement: Natural conversation increases usage
- Cost Reduction: Automated support reduces operational costs

#### **Pilot Plan:**

1. Phase 1 (Months 1-2): Internal testing, bug fixes, performance optimization
2. Phase 2 (Months 3-4): Beta testing with 100-500 users
3. Phase 3 (Months 5-6): Gradual rollout to 10% of user base
4. Phase 4 (Months 7-12): Full rollout with monitoring and improvements

#### **ROI Projection:**

- Year 1: 30% reduction in call center queries
- Year 2: 50% reduction in call center queries
- Year 3: 70% reduction in call center queries
- Cost Savings: \$500K - \$2M annually (depending on scale)

## **8.4 Compliance & Privacy**

#### **Regulatory Compliance:**

- RBI guidelines for UPI payments
- PIN security requirements
- User consent mechanisms
- Audit trail and logging
- Data privacy (embeddings only, not raw audio)

#### **Privacy Features:**

- Voice samples not stored
- PII detection and prevention
- Encrypted data storage (production)
- User consent for data usage

## 9. Repository Structure

### 9.1 Code Organization

...

```
vaani-banking-voice-assistant/
├── frontend/          React frontend (Port 5173)
│   ├── src/
│   │   ├── pages/      Page components (Login, Chat, Profile, Transactions)
│   │   ├── hooks/      Custom hooks (useVoiceRecognition, useTextToSpeech)
│   │   ├── services/    API clients (aiClient, api)
│   │   └── components/ Reusable components
│   └── package.json
└── backend/           FastAPI backend (Port 8000)
    ├── api/            API routes, schemas, security
    ├── db/             Database models, repositories, services
    │   ├── models/      SQLAlchemy models
    │   ├── repositories/ Data access layer
    │   └── services/    Business logic
    └── app.py          FastAPI application
└── ai/               AI backend (Port 8001)
```

```
agents/      LangGraph agents
    banking_agent.py
    upi_agent.py
    rag_agents/  Loan, Investment, Support agents
orchestrator/ HybridSupervisor, IntentRouter
services/     LLM, RAG, Guardrail services
tools/        Banking and UPI tools
chroma_db/   Vector databases
main.py       FastAPI server

documentation/ Comprehensive documentation
    system_architecture.md
    ai_architecture.md
    backend_architecture.md
    guardrails_implementation.md
    hindi_support.md
    upi_payment_flow.md

requirements.txt Python dependencies
...
```

## 11. Conclusion

Vaani Banking Voice Assistant represents a comprehensive, production-ready solution for voice-enabled financial operations. With its innovative multi-agent architecture, bilingual RAG system, comprehensive security guardrails, and RBI-compliant UPI implementation, Vaani demonstrates:

1. Deep Technical Understanding: Modern microservices architecture, LangGraph orchestration, vector databases
2. Security & Compliance: Voice biometrics, guardrails, RBI compliance, privacy-first design
3. Innovation: Hybrid supervisor pattern, bilingual RAG, voice-first banking
4. Business Impact: Complete banking operations, accessibility, scalability, ROI potential
5. Production Readiness: Observability, error handling, scalability, comprehensive documentation

Document Version: 1.0

Last Updated: Nov 2025

Prepared By: Vaani Development Team